MASTER SYLLABUS · VERSION 1.1 (AS BUILT) · 12 OCTOBER 2026

Agentic AI Engineering

From zero to production-grade, governed AI agents — a 30-week programme with an India-first lens

Prepared for Arunkumar Maniam Rajan · Chennai, India (IST)

Version 1.1 updates the plan to what was actually built and run (sections 11–13)

Tool versions verified on PyPI and official docs on 11 Oct 2026

Reference benchmark: IITM Pravartak × Emeritus Professional Certificate Programme in Agentic AI and Applications

Contents

1. Programme overview **3**

1.1 Design principles 3

2. Audience and prerequisites **4**

2.1 Intended learner 4

2.2 Prerequisites 4

3. How this programme improves on the reference **4**

4. Module list **5**

5. Week-by-week schedule **7**

6. Learning outcomes per module **8**

7. Assessment plan **11**

7.1 Capstone rubric (all capstones) 11

7.2 Capstone briefs 11

8. Tools and setup **12**

8.1 Software stack (versions verified 11 Oct 2026) 12

8.2 Hardware and accounts 13

8.3 Repository layout 13

9. Learning approach and policies **13**

10. Reference programme and sources **14**

11. As-built deliverables **15**

12. What ran on what backend **16**

13. Known limitations **17**

Appendix A. Week 0 — Python bridge for Kotlin/Java developers (optional) **18**

Appendix B. Glossary **18**

1. Programme overview

Agentic AI Engineering is a 30-week, project-driven programme that takes an experienced software engineer from first principles of large language models (LLMs) to designing, evaluating, securing and operating multi-agent systems in production. It is organised in five phases that deliberately mirror — and extend — the five pillars of the IITM Pravartak × Emeritus *Professional Certificate Programme in Agentic AI and Applications* used as the reference benchmark (a 5-month core plus a 6-week specialisation).

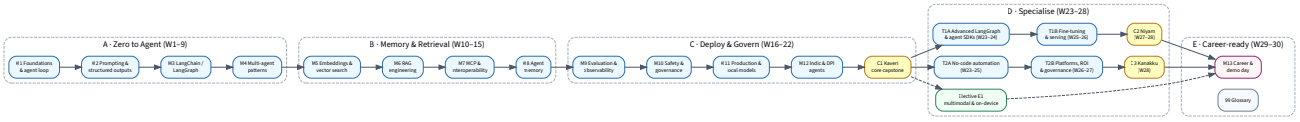


Figure 1. Programme map: five phases, thirteen modules, three capstones (C1 core; C2 for Track 1 or C3 for Track 2), two specialisation tracks split into parts A and B, and the optional elective E1.

> **Version 1.1: as built**
>
> Version 1.0 of this syllabus was the plan. Version 1.1, dated 12 October 2026, records what was actually built: 23 PDFs (section 11), which backend each lab ran on (section 12), and the known limitations (section 13). The capstones are named **C1** (core, Kaveri), **C2** (Track 1, Niyam) and **C3** (Track 2, Kanakku). Each track is delivered in two parts: **T1A/T1B** and **T2A/T2B**. Where the plan and the build differ, sections 11–13 take precedence.

Duration	30 weeks: 22-week core (≈5 months) + 6-week specialisation + 2-week career sprint
Effort	10–12 hours/week (≈ 330 hours total): ~2 h reading, ~1 h quiz/review, 6–8 h labs & projects, 1 h discussion
Format	Self-paced textbook-style module PDFs, runnable lab repositories, weekly checkpoints, two capstones per learner (C1 plus C2 or C3)
Language	Python 3.12/3.13 (labs tested on 3.13.5); optional Kotlin/Android elective
Cost to run	Every lab runs offline on a laptop (mock LLM or local small model). A cloud API key is optional, and none was used to build the course
Outcome	A portfolio of three production-quality agent repositories, an evaluation & safety dossier, and interview readiness

1.1 Design principles

1. **First principles before frameworks.** You build an agent loop by hand (Module 1) before using LangChain, LangGraph, CrewAI or vendor SDKs, so frameworks become conveniences rather than magic.
2. **Evaluation is a first-class skill.** Every module from M1 ships a small evaluation harness; M9 makes it rigorous.
3. **Safety by construction.** Approval gates for side-effecting tools appear in Module 1; prompt-injection defence, least privilege and governance are graded deliverables, not slides.
4. **Framework-plural, concept-stable.** Concepts (state, tools, memory, handoffs, retrieval, evaluation) are taught once and mapped onto several frameworks, because the ecosystem moves monthly (e.g., AutoGen entered maintenance mode and Microsoft Agent Framework became its successor).

5. **Runs without API keys.** Mock LLMs and local small models (Ollama, Hugging Face on CPU) make every lab reproducible, cheap and private; cloud models are a drop-in upgrade.
6. **India-first context.** Indic languages, token-cost effects of scripts, UPI/ONDC/Account-Aggregator-style workflows (against mocks and sandboxes only), DPDP Act 2023 compliance and INR cost models are woven throughout.
7. **Built for this learner.** Analogies and an elective draw on Android and multimedia engineering (coroutines, lifecycles, streaming pipelines, on-device inference).

2. Audience and prerequisites

2.1 Intended learner

Professional software engineers (3+ years) who want to become AI/agent engineers — in particular mobile and multimedia engineers who already think in terms of state machines, asynchronous pipelines, performance budgets and device constraints. Track 2 additionally serves tech leads and architects who will own agent strategy.

2.2 Prerequisites

Area	Required	Helpful (taught as needed)
Programming	Comfortable in one typed language (Kotlin/Java/C++); basic Python syntax	Python typing, asyncio, packaging (taught in M1)
Maths	High-school algebra; vectors & matrices; basic probability	Softmax, dot products, entropy (taught in M1)
Tools	Git, a terminal, HTTP/JSON APIs	Docker (taught in M11)
Hardware	Any 64-bit laptop with 8 GB RAM (16 GB recommended)	A GPU is **not** required; Track 1 fine-tuning uses free/low-cost cloud GPUs

Self-check before Week 1: write a Python function that fetches JSON from a URL and validates two fields; explain what a matrix multiplication of shapes (4×8)·(8×3) produces. If either is hard, do the optional *Week 0 Python bridge* (Appendix A).

3. How this programme improves on the reference

The reference programme's syllabus is strong on breadth. This design keeps all of its topics and adds the following:

Gap / opportunity in the reference outline	What this programme adds	Where
Frameworks introduced alongside fundamentals	Hand-built agent loop, typed client layer and unit tests before any framework	M1
Structured outputs & tool calling implicit	A dedicated module: JSON Schema, Pydantic contracts, strict/native tool calling, constrained decoding, validation-retry	M2

Gap / opportunity in the reference outline	What this programme adds	Where
Evaluation focused on hallucination detection	Full evaluation discipline: golden sets, trajectory evals, LLM-as-judge calibration, pass@k, CI regression gates, OpenTelemetry GenAI conventions	M1 (seed), M9
Cost and latency not explicit	Token economics in INR, prompt/semantic caching, model cascades and routing, TTFT/TPOT budgets, quantisation	M1, M11
Local/small models absent	Ollama, llama.cpp, vLLM, GGUF quantisation, SLM selection, offline labs from day one	M1, M11, T1
Agent safety mostly at the end	Approval gates from M1; red-team suite with measured attack-success rate; dual-LLM/privilege-separation patterns	M1, M10
Generic use cases	Indic NLP (AI4Bharat, Bhashini, Sarvam), Indic token fertility, UPI/ONDC/AA/DigiLocker-style workflows on mocks, DPDP Act	M12 + threaded
Framework landscape as of 2025	Updated to Oct 2026: LangChain 1.x/LangGraph 1.2, Microsoft Agent Framework (AutoGen successor), MCP 2026-07-28 spec with extensions, A2A	M3, M4, M7
Memory treated with retrieval	Separate memory-design module with consent, retention and forgetting policies	M8
No on-device/multimodal path	Elective E1: Gemini Nano / ML Kit GenAI, LiteRT-LM (MediaPipe LLM Inference is now maintenance-only), ExecuTorch, ONNX Runtime, llama.cpp and MLC LLM on Android	E1
Career support generic	Role-mapped portfolio (3 repos), a 140-question interview bank with model answers, worked system-design interviews, Android→AI transition narrative	M13

4. Module list

#	Module	Weeks	Phase	One-line description
M1	Foundations: Python for AI Engineering, LLM Fundamentals & the Agent Loop	1–3	A	Typed, async, tested Python; transformers from scratch; build a tool-using agent with no framework
M2	Prompting, Structured Outputs & Tool Calling	4–5	A	Output contracts with JSON Schema/Pydantic, native vs prompted tools, constrained decoding
M3	Frameworks I: LangChain 1.x & LangGraph	6–7	A	Stateful graphs, checkpointing, human-in-the-loop interrupts, streaming
M4	Multi-Agent Systems & Design Patterns	8–9	A	CrewAI, Microsoft Agent Framework / AutoGen Studio, AG2; supervisor, handoff, debate, plan-execute
M5	Embeddings & Vector Search	10–11	B	ChromaDB & FAISS, ANN index families, hybrid BM25 + vector, cross-encoder re-ranking
M6	RAG Engineering	12–13	B	Ingestion, chunking, query rewriting, agentic & graph RAG, citations, RAG evaluation
M7	Model Context Protocol & Interoperability	14	B	MCP (spec 2026-07-28) servers/clients, transports, auth, security; A2A
M8	Agent Memory Design	15	B	Working/episodic/semantic/procedural memory, write & forget policies, context engineering
M9	Evaluation, Observability & Hallucination Detection	16–17	C	LangSmith/Langfuse tracing, eval datasets, LLM-as-judge, CI gates

#	Module	Weeks	Phase	One-line description
M10	Agent Safety, Security & AI Governance	18	C	OWASP LLM Top 10, prompt-injection defence, guardrails, DPDP Act, EU AI Act, ISO/IEC 42001
M11	Production Engineering, Cost/Latency & Local Models	19	C	FastAPI + Docker, streaming, caching, routing, quantised local models
M12	Indic & Digital-Public-Infrastructure Agents	20	C	Voice-first Indic agents; UPI/ONDC/AA-style workflows against sandboxes
C1	Core capstone: Kaveri, a bilingual customer-operations agent	21–22	C	Integrates M1–M12 under a graded rubric
T1A	Track 1 Part A — Advanced LangGraph and agent SDKs	23–24	D	Durable graphs, OpenAI Agents SDK, Google ADK, Microsoft Agent Framework, A2A
T1B	Track 1 Part B — Fine-tuning and serving	25–26	D	LoRA on CPU (QLoRA documented), evaluation of fine-tunes, vLLM/Kubernetes serving (manifests validated)
C2	Track 1 capstone: Niyam, a multi-agent regulatory-compliance RAG	27–28	D	LoRA router, verifier, eval gates, load test
T2A	Track 2 Part A — No-code automation and process discovery	23–25	D	n8n (self-hosted), Zapier and Make (recipes), opportunity mapping
T2B	Track 2 Part B — Enterprise platforms, ROI and governance	26–27	D	Platform scorecard, INR ROI with Monte Carlo, DPDP governance pack
C3	Track 2 capstone: Kanakku, invoice-to-pay for an Indian SME	28	D	n8n + local model, GST validation, approvals, Tally mock, ROI and governance
E1	Elective — Multimodal & On-device Agents (Android)	parallel	D	Gemini Nano/ML Kit GenAI, LiteRT-LM, ExecuTorch, llama.cpp; small VLMs, ASR; hybrid device–cloud routing
M13	Career & Certification	29–30	E	Portfolio, resume, LinkedIn, 140-question interview bank, certification plan, demo day
—	Course glossary	—	—	`99-glossary.pdf` : every module's glossary merged, plus acronyms

5. Week-by-week schedule

Each week lists the theory focus, the hands-on lab and the checkpoint. “Q” = module quiz, “MP” = mini-project.

Wk	Module	Theory focus	Lab / build	Checkpoint
0	(optional)	Python bridge for Kotlin/Java developers	Kotlin→Python kata set	Self-check
1	M1	What agents are; Python for AI engineering (uv, typing, Pydantic, asyncio, httpx, pytest); AI-assisted development	Lab 1.1 env & provider-agnostic client; Lab 1.7 async + retries	Env verified
2	M1	Tokens, BPE, embeddings, self-attention, transformer blocks, training pipeline, decoding	Labs 1.2–1.5: BPE, sampling, attention from scratch, tiny LLM on CPU	—
3	M1	Calling LLMs, tool calling, the ReAct loop, first evals, approval gates	Lab 1.6 agent loop + mini-eval across mock/local backends	Q1, labs
4	M2	Prompt anatomy, few-shot, CoT trade-offs, prompt versioning	Prompt test bench with regression cases	—
5	M2	JSON Schema, Pydantic contracts, native/strict tool calling, constrained decoding, instructor, pydantic-ai	GST invoice extractor (EN/TA/HI)	Q2, MP1
6	M3	LangChain 1.x agents & middleware; LangGraph StateGraph, reducers, edges	Port the M1 agent to LangGraph	—
7	M3	Checkpointers, interrupts (HITL), streaming, durable execution, Studio	Support agent with approval + persistence	Q3
8	M4	Workflow vs agent patterns; supervisor, hierarchical, handoff/swarm, debate, reflection	CrewAI crew & flow	—
9	M4	Microsoft Agent Framework workflows, AutoGen Studio prototyping, AG2; A2A intro	Same task in 3 frameworks; cost/quality report	Q4, MP2
10	M5	Embedding models (incl. multilingual/Indic), metrics, exact vs ANN	FAISS Flat/IVF/HNSW benchmark	—
11	M5	ChromaDB, metadata filters, BM25, hybrid fusion (RRF), cross-encoder re-ranking	Bilingual hybrid search + rerank	Q5
12	M6	RAG architecture, parsing, chunking strategies, contextual retrieval	Chunking ablation study	—
13	M6	Query rewriting, HyDE, agentic/corrective RAG, GraphRAG, citations, ragas/deepeval	Government-circulars RAG with citations	Q6, MP3
14	M7	MCP architecture, primitives, transports, auth, extensions, security; A2A	MCP server for a mock ONDC catalogue	Q7
15	M8	Memory taxonomy, stores, write/forget policies, consent, context engineering	Memory-enabled finance assistant	Q8
16	M9	Eval taxonomy, golden sets, metrics, LLM-as-judge, trajectory evals	100-item eval set for the RAG agent	—
17	M9	Tracing (LangSmith, Langfuse, OpenTelemetry GenAI), dashboards, CI gates, hallucination detection	Instrumented agent + CI regression gate	Q9

Wk	Mod-ule	Theory focus	Lab / build	Checkpoint
18	M10	OWASP LLM Top 10, prompt injection, guardrails, sandbox-ing, DPDP/EU AI Act/NIST/ISO 42001	Red-team & defend; measure attack success rate	Q10
19	M11	FastAPI, Docker, streaming, queues; cost/latency engineer-ing; local & quantised models	Containerised agent with model cas-cade & ₹ cost report	Q11
20	M12	Indic NLP stack, ASR/TTS, code-mixing; UPI/ONDC/AA/Digi-Locker patterns	Tamil voice-to-action agent on mocks	Q12
21	C1	Capstone design review: architecture, threat model, eval plan	Build sprint 1	Design doc
22	C1	Hardening, evaluation, documentation	Build sprint 2 + demo video	**Core capstone**
23	T1A / T2A	T1: advanced LangGraph · T2: n8n, Zapier, Make	Track lab 1	—
24	T1A / T2A	T1: OpenAI Agents SDK & Google ADK · T2: workflows that survive failure, approvals	Track lab 2	T1 quiz A
25	T1B / T2A	T1: LoRA/QLoRA with PEFT+TRL · T2: process discovery & opportunity mapping	Track lab 3	T2 quiz A
26	T1B / T2B	T1: serving & scaling (vLLM, K8s) · T2: enterprise platforms; ROI & business case in INR	Track lab 4	T1 quiz B
27	C2 / T2B	T1: C2 Niyam build · T2: governance operating model	C2 sprint 1 · T2B lab	Design review · T2 quiz B
28	C2 / C3	Specialisation capstone (C2 hardening, or the C3 Kanakku week)	Capstone sprint	**Specialisation capstone**
29	M13	Portfolio, resume, LinkedIn	Portfolio repos polished	Peer review
30	M13	Interview prep: concepts, agent system design, take-home	Mock interviews + demo day	**Final defence, completion**

6. Learning outcomes per module

Outcomes are written to be observable and assessable. On completing each module you will be able to:

M1 — Foundations. (1) set up a reproducible, pinned Python environment and structure an AI project with typed models, async I/O, retries and tests; (2) explain tokenization, embeddings, self-attention, the de-coder-only transformer, KV caching and the pre-training → SFT → preference-optimisation pipeline; (3) im-plement softmax, temperature, top-k and top-p sampling and predict their effect; (4) call LLMs through a provider-agnostic client (mock, local, cloud); (5) build, trace and unit-test a tool-using agent loop with an approval gate; (6) quantify how script choice affects token counts and therefore cost for Indic languages.

M2 — Structured outputs & tool calling. (1) design output contracts with JSON Schema and Pydantic; (2) choose between native, strict and prompted tool calling and constrained decoding; (3) design tools with good names, granularity, idempotency and error messages; (4) build validation-retry loops that reach ≥95% schema validity.

M3 — LangChain & LangGraph. (1) model an agent as a typed state graph with reducers and conditional edges; (2) persist and resume conversations with checkpointers; (3) add human-in-the-loop interrupts; (4) stream tokens and intermediate steps; (5) decide when a framework adds value over hand-written code.

M4 — Multi-agent systems. (1) select an appropriate pattern (chain, router, orchestrator-workers, evaluator-optimiser, supervisor, handoff, debate) for a problem; (2) implement it in CrewAI, LangGraph and Microsoft Agent Framework; (3) measure the cost, latency and quality trade-offs of adding agents; (4) explain A2A's role.

M5 — Embeddings & vector search. (1) choose and evaluate embedding models including multilingual ones; (2) build and tune FAISS (Flat/IVF/PQ/HNSW) and ChromaDB indexes; (3) implement hybrid search with RRF and cross-encoder re-ranking; (4) report recall@k, MRR and nDCG.

M6 — RAG engineering. (1) design ingestion and chunking for heterogeneous documents incl. Indic scripts; (2) implement query rewriting, multi-query, HyDE and agentic/corrective RAG; (3) enforce citations; (4) evaluate faithfulness and context precision/recall and run ablations.

M7 — MCP & interoperability. (1) explain MCP's host/client/server model, primitives and transports per the 2026-07-28 specification; (2) build and secure an MCP server with the Python SDK/FastMCP; (3) consume it from two agent frameworks; (4) identify tool-poisoning and confused-deputy risks.

M8 — Agent memory. (1) classify memory types and select stores; (2) design write, retrieval and forgetting policies; (3) implement consent-aware long-term memory; (4) apply context-engineering techniques (compaction, scratchpads, sub-agents).

M9 — Evaluation & observability. (1) build golden datasets and trajectory evals; (2) calibrate LLM-as-judge against human labels; (3) instrument agents with tracing (LangSmith/Langfuse/OpenTelemetry); (4) detect hallucinations with entailment and citation checks; (5) gate releases in CI.

M10 — Safety, security & governance. (1) threat-model an agent with the OWASP LLM Top 10; (2) implement layered prompt-injection defences and measure attack success rate; (3) configure guardrails and PII redaction for Indian identifiers; (4) map obligations under the DPDP Act 2023, EU AI Act and NIST AI RMF; (5) write an agent card and incident runbook.

M11 — Production, cost & local models. (1) ship an agent as a FastAPI + Docker service with streaming; (2) build a cost model in INR and reduce it with caching, routing and cascades; (3) meet TTFT/TPOT budgets; (4) select, quantise and serve small local models with Ollama/llama.cpp/vLLM.

M12 — Indic & DPI agents. (1) assemble an Indic voice pipeline (ASR → LLM → TTS) with translation/transliteration; (2) design agents that assist UPI/ONDC/AA-style flows without bypassing regulated consent steps; (3) evaluate multilingual quality and fairness; (4) design for low bandwidth and accessibility.

C1 — Core capstone. Integrate tools, MCP, RAG, memory, approval gates, tracing, evaluation and security into one deployable bilingual agent, with documented trade-offs.

T1 — Advanced engineering. (1) build durable multi-agent graphs; (2) compare OpenAI Agents SDK, Google ADK and Microsoft Agent Framework on one problem; (3) fine-tune a small model with LoRA/QLoRA for tool routing and evaluate it; (4) serve and load-test at scale.

T2 — Strategy & enterprise workflows. (1) build agentic automations in n8n, Zapier and Make; (2) evaluate enterprise agent platforms for fit, lock-in and data residency; (3) build an ROI model in INR; (4) design an AI governance operating model and change plan.

E1 — Multimodal & on-device. (1) choose among Gemini Nano (ML Kit GenAI), LiteRT-LM, ExecuTorch, ONNX Runtime and llama.cpp for an Android feature; (2) measure tokens/s and memory against a budget; (3) make small-VLM and ASR output safe with validators; (4) design hybrid on-device/cloud routing within privacy, latency, memory and battery budgets.

M13 — Career. (1) present three portfolio repositories with eval reports and demos; (2) answer conceptual and agent system-design interview questions; (3) articulate a credible Android→AI-engineer transition story; (4) write resume bullets traceable to logged results; (5) plan the next 90 days against a skills matrix.

7. Assessment plan

Component	Count	Weight	Notes
Labs	12 module lab sets (M1–M12)	30% (2.5% each)	Auto-graded tests + short reflection; must include run log
Quizzes	12 (M1–M12)	15% (1.25% each)	20 questions; two attempts; answer key released after deadline
Mini-projects	MP1 (M2), MP2 (M4), MP3 (M6)	15% (5% each)	Graded with a 4-criterion rubric
Core capstone (C1)	1	20%	Rubric below; live demo + code review
Specialisation capstone	1 (Track 1 or 2)	15%	Track-specific rubric
Career portfolio & participation	1	5%	Portfolio, resume, mock interview
Total		**100%**	Pass ≥ 60% overall **and** ≥ 50% in each capstone **and** the final defence passed; distinction ≥ 85%

Track quizzes (T1A, T1B, T2A, T2B) and the E1 quiz are formative (0%). The **final defence** (60 minutes: capstone demo, technical questions, an unseen system design) is a pass/fail gate. Certificates: Certificate of Completion (Track 1 or Track 2), with Distinction at ≥ 85%, an Elective E1 endorsement if the E1 project scores ≥ 60/100, and a Statement of Participation otherwise (labs for at least 6 modules submitted). It is a course certificate, not an accredited award. Rules and tests: M13 §13.7 and Lab 13.3.

7.1 Capstone rubric (all capstones)

Criterion	Weight	Excellent (90–100%) looks like
Functionality	30%	All required capabilities work end-to-end on the demo scenarios and the hidden test set
Engineering quality	20%	Typed, tested (≥70% coverage on core logic), containerised, clear configuration, pinned dependencies
Evaluation	20%	≥50-case golden set, task & trajectory metrics, judge calibration, regression gate, honest error analysis
Safety & governance	15%	Threat model, injection suite with attack-success rate, approval gates, PII handling, DPDP mapping
Documentation & demo	15%	Architecture diagram, ADRs, cost/latency budget in INR, ≤ 3-minute demo video

7.2 Capstone briefs

C1 core capstone — "Kaveri" bilingual operations agent. A customer-operations agent for an Indian D2C retailer or fintech that answers policy questions (RAG with citations), checks order/payment status (mock UPI/ONDC tools via an MCP server), remembers customer preferences with consent, escalates refunds through human approval, and works in English plus Tamil or Hindi. Delivered as a FastAPI + Docker service with tracing and evaluation reports.

C2 Track 1 capstone — "Niyam" production multi-agent RAG system. For example, a regulatory-compliance copilot over public RBI/SEBI circulars: planner, retriever, analyst and verifier agents; a LoRA-fine-tuned small router model; hybrid retrieval with re-ranking; full observability; load test; eval gates.

C3 Track 2 capstone — "Kanakku" end-to-end enterprise agentic workflow. For example, invoice-to-pay for an Indian SME: email/WhatsApp intake, GST invoice extraction and validation, three-way match against a mock ERP, approval routing, exception handling by an agent, built with n8n/Make plus code where needed; delivered with an ROI model, risk register and governance pack.

8. Tools and setup

8.1 Software stack (versions verified 11 Oct 2026)

Purpose	Tool	Version	Used in
Language	Python	3.13 (3.12 also fine; some libraries require < 3.14)	all
Env & packaging	uv or venv + pip	latest	all
Data validation	pydantic	2.14.0	M1+
HTTP	httpx	0.28.1	M1+
Local models	Ollama	0.40.3	M1, M11
Local models (Python)	transformers / torch (CPU)	5.19.0 / 2.14.1	M1, T1
Agent framework	langchain / langchain-core	1.4.7 / 1.7.3	M3+
Graph orchestration	langgraph	1.2.19	M3+
Multi-agent	crewai	1.15.34	M4
Multi-agent	agent-framework (Microsoft)	1.21.0	M4, T1
Multi-agent (legacy/maintenance)	autogen-agentchat / autogenstudio	0.7.5 / 0.4.2.2	M4
Multi-agent (community)	ag2	1.1.2	M4
Vendor SDKs	openai-agents / google-adk	0.23.1 / 2.11.0	T1
MCP	mcp (Python SDK) / fastmcp; spec revision	2.3.0 / 4.1.5; 2026-07-28	M7+
Vector search	chromadb / faiss-cpu	1.5.9 / 1.15.1	M5+
Embeddings & re-ranking	sentence-transformers	6.1.0	M5+
Structured outputs	instructor / pydantic-ai	1.17.0 / 2.55.0	M2
Evaluation	ragas / deepeval	0.4.3 / 4.2.8	M6, M9
Observability	langsmith / langfuse	0.14.7 / 4.17.0	M9+
Serving	fastapi / uvicorn	0.143.0 / 0.54.0	M11+
Fine-tuning	peft / trl	0.21.2 / 1.15.0	T1
Gateway (optional)	litellm	1.105.0	M11
No-code (Track 2)	n8n, Zapier, Make	SaaS / self-hosted	T2

Version policy

Versions above were read from PyPI on 11 October 2026. Each module pins its own `requirements.txt` and is re-verified when the module is built. When a framework ships a breaking change mid-course, the module notes the delta rather than silently changing code.

8.2 Hardware and accounts

- **Minimum:** 8 GB RAM laptop, 20 GB free disk. **Recommended:** 16 GB RAM. No GPU needed for core modules.
- **Optional accounts:** one cloud LLM provider (any OpenAI-compatible endpoint), LangSmith or self-hosted Langfuse, GitHub, Docker Hub. Track 1's QLoRA and vLLM steps need a GPU (e.g., a single 24 GB card for ~6–10 hours); as built, they are documented and validated but were not run (section 13), and the LoRA fine-tune ran on the CPU.
- **Budget guidance:** with local models for development and a cloud model only for final evaluations, total API spend for the whole programme can be kept under ₹3,000; Track 1 GPU time is the main additional cost.

8.3 Repository layout

```
ai-course/
  README.md                   # index of every PDF with page counts (generated)
  00-master-syllabus.pdf … 13-module-13-career.pdf, C1/C2/C3, T1A/T1B/T2A/T2B, E1, 99-glossary.pdf
  course-plan.md              # design contract: specs, template, style rules, version pins
  labs/module-01 … module-13/ # per-module labs, tests, run_all.sh, outputs/logs/, outputs/RUN_LOG.md
  labs/capstone-core|capstone-t1|capstone-t2/   # C1, C2, C3
  labs/track1-a|track1-b|track2-a|track2-b/     # track parts
  labs/elective-e1/           # E1 Python labs and Kotlin samples (not compiled)
  labs/course-glossary/ labs/course-index/      # generators for 99-glossary.pdf, README and section
11
  build/                      # build.py, style.css, diagrams/*.dot, check_pdf.sh
  src/                        # Markdown sources (src/parts/ per chapter)
  research/                   # version snapshots and link-check results
```

9. Learning approach and policies

- **Weekly rhythm:** Mon–Tue read theory; Wed–Fri labs; Sat checkpoint (quiz or deliverable); Sun rest/review.
- **AI-assistant policy:** coding assistants are encouraged (M1 teaches a disciplined workflow), but every submission must include tests you understand and a short "what I verified" note. Quizzes are closed-assistant.
- **Reproducibility:** submissions must run with `./run_all.sh mock` on a clean machine.
- **Ethics & law:** never test against real payment, banking or government systems; use mocks or official sandboxes; never commit secrets or personal data.

10. Reference programme and sources

- Reference programme: IITM Pravartak × Emeritus, *Professional Certificate Programme in Agentic AI and Applications* — https://iitmpravartak.emeritus.org/professional-certificate-programme-in-agentic-ai-and-applications (the URL `iitmpravartak.emeritus.com` given in the brief returned HTTP 404 on 11 Oct 2026; the `.org` page above is the live programme page).
- AutoGen maintenance-mode notice and Microsoft Agent Framework migration guide — https://github.com/microsoft/autogen, https://learn.microsoft.com/en-us/agent-framework/migration-guide/from-autogen/
- MCP specification, revision 2026-07-28 — https://modelcontextprotocol.io/specification/2026-07-28
- Anthropic, *Building effective agents* — https://www.anthropic.com/engineering/building-effective-agents
- OWASP GenAI Security Project, LLM Top 10 — https://genai.owasp.org/llm-top-10/
- Package versions — PyPI project pages, e.g. https://pypi.org/project/langgraph/1.2.19/ (full list in `research/`).

11. As-built deliverables

The course was delivered as 23 PDFs. The table below is generated by `labs/course-index/inventory.py` from `pdfinfo` and the files on disk (12 October 2026), so the page counts are the real ones. The same table opens `README.md`.

File	Contents	Weeks	Pages	Size	Lab code
`00-master-syllabus.pdf`	Master syllabus (as built, v1.1)	—	18	174 KB	—
`01-module-1.pdf`	M1 Foundations: Python, LLM fundamentals, the agent loop	1–3	63	649 KB	`labs/module-01`
`02-module-2.pdf`	M2 Prompting, structured outputs and tool calling	4–5	56	492 KB	`labs/module-02`
`03-module-3.pdf`	M3 LangChain 1.x and LangGraph	6–7	52	389 KB	`labs/module-03`
`04-module-4.pdf`	M4 Multi-agent systems and design patterns	8–9	44	419 KB	`labs/module-04`
`05-module-5.pdf`	M5 Embeddings and vector search	10–11	64	762 KB	`labs/module-05`
`06-module-6.pdf`	M6 RAG engineering	12–13	84	744 KB	`labs/module-06`
`07-module-7.pdf`	M7 Model Context Protocol and interoperability	14	44	336 KB	`labs/module-07`
`08-module-8.pdf`	M8 Agent memory design	15	43	338 KB	`labs/module-08`
`09-module-9.pdf`	M9 Evaluation, observability and hallucination detection	16–17	56	463 KB	`labs/module-09`
`10-module-10.pdf`	M10 Agent safety, security and AI governance	18	52	404 KB	`labs/module-10`
`11-module-11.pdf`	M11 Production engineering, cost/latency and local models	19	45	371 KB	`labs/module-11`
`12-module-12.pdf`	M12 Indic and digital-public-infrastructure agents	20	50	395 KB	`labs/module-12`
`C1-core-capstone.pdf`	C1 Core capstone: Kaveri bilingual operations agent	21–22	33	282 KB	`labs/capstone-core`
`T1A-track1-advanced-engineering.pdf`	T1A Advanced LangGraph and agent SDKs	23–24	51	353 KB	`labs/track1-a`
`T1B-track1-finetuning-serving.pdf`	T1B Fine-tuning (LoRA) and serving	25–26	51	377 KB	`labs/track1-b`
`C2-T1-capstone.pdf`	C2 Track 1 capstone: Niyam multi-agent regulatory RAG	27–28	37	319 KB	`labs/capstone-t1`
`T2A-track2-nocode-automation.pdf`	T2A No-code/low-code automation and process discovery	23–25	59	433 KB	`labs/track2-a`

File	Contents	Weeks	Pages	Size	Lab code
`T2B-track2-enterprise-roi-governance.pdf`	T2B Enterprise platforms, ROI and governance	26–27	50	449 KB	`labs/track2-b`
`C3-T2-capstone.pdf`	C3 Track 2 capstone: Kanakku invoice-to-pay	28	37	403 KB	`labs/capstone-t2`
`E1-multimodal-on-device-android.pdf`	E1 Multimodal and on-device agents on Android (optional)	parallel to 23–28	49	392 KB	`labs/elective-e1`
`13-module-13-career.pdf`	M13 Career and certification	29–30	55	348 KB	`labs/module-13`
`99-glossary.pdf`	Course glossary and acronyms	—	20	122 KB	`labs/course-glossary`
Total	23 PDFs	30 weeks	**1113**		

Changes from version 1.0 of the plan:

- **Capstone naming.** The three capstones are C1 (core, Kaveri), C2 (Track 1, Niyam) and C3 (Track 2, Kanakku). Each learner completes C1 and one of C2 or C3.
- **Track splits.** Each track was too large for one booklet, so each is split into part A and part B. Track 1: T1A (weeks 23–24) and T1B (weeks 25–26), then C2 (weeks 27–28). Track 2: T2A (weeks 23–25) and T2B (weeks 26–27), then C3 (week 28).
- **E1.** The elective is a separate optional PDF that can run alongside either track. Because the Google AI Edge documentation now marks MediaPipe LLM Inference as maintenance-only, E1 teaches LiteRT-LM and ML Kit GenAI instead.
- **Glossary.** Appendix B now points to `99-glossary.pdf`. It has 336 unique terms (merged from 397 module entries) and 69 acronyms.

12. What ran on what backend

No cloud LLM API was called anywhere in the course. “Mock” means a deterministic scripted model, so results can be reproduced exactly. “Ollama” means a small Qwen2.5-Instruct model (0.5B, 1.5B or 3B, Q4_K_M) served locally by Ollama 0.40.3. Every run was on one 8-core x86-64 CPU-only VM with 15 GB of RAM and Python 3.13.5. Each module’s Appendix A has the full run log.

Unit	Backends actually used	Notable items not run
M1	Mock; Ollama qwen2.5 0.5B; Hugging Face SmolLM2-360M on CPU	—
M2	Mock; Ollama qwen2.5 1.5B; HF	—
M3	Mock; Ollama qwen2.5 1.5B (LangGraph)	—
M4	Mock; Ollama 1.5B (including CrewAI in its own venv)	—
M5	HF embedding models on CPU; FAISS/HNSW; no generator	—
M6	HF retrievers and rerankers; mock; Ollama 1.5B/3B	some Ollama variants; Ragas/DeepEval (code only)
M7	Mock; a short Ollama 1.5B run; local MCP servers	—

Unit	Backends actually used	Notable items not run
M8	Deterministic memory labs; Ollama 1.5B (Lab 8.3)	—
M9	Ollama 1.5B; NLI models on CPU; self-hosted Arize Phoenix	hosted tracing services (no accounts)
M10	SusceptibleMock; Ollama 1.5B; offline classifiers; benign canary attacks only	NeMo Guardrails, Guardrails AI, LLM Guard, garak, promptfoo, Prompt Guard 2 (gated)
M11	Mock; Ollama 0.5B/1.5B/3B; rootless podman	podman compose, locust, vLLM (no GPU), speculative decoding
M12	Offline tokenisers (Qwen, GPT-2, Sarvam and others); Ollama 0.5B/1.5B/3B	every real DPI/Indic API, ASR/TTS, IndicTrans2, Indic LLM inference
C1	Mock; Ollama 1.5B	—
T1A	Mock; Ollama 1.5B through the OpenAI Agents SDK, Google ADK and LangGraph	hosted model APIs
T1B	CPU LoRA fine-tune of Qwen2.5-0.5B (PEFT/TRL); kubeconform/kube-linter manifest checks	QLoRA/bitsandbytes and vLLM (GPU only); a real Kubernetes cluster
C2	HF LoRA router; Ollama 1.5B; podman	kind cluster (failed to create under rootless podman)
T2A	n8n 2.42.6 in podman; Ollama 1.5B	Zapier and Make (recipes only, no accounts)
T2B	Python ROI model, Monte Carlo, xlsx; pricing read from official pages	every enterprise agent platform (no accounts)
C3	n8n plus Ollama 1.5B; a rules-only reference pipeline	real Tally, bank or e-mail integrations (all mocked)
E1	HF SmolVLM, Whisper and MMS on CPU; Ollama as an on-device proxy	Kotlin samples not compiled; no Android device or emulator; Gemini Nano
M13	Python standard library only (results bank, question bank, grader, portfolio check)	—
Glossary	Python standard library only	—

13. Known limitations

1. **No cloud API key.** All LLM results come from mocks or small local models. Quality numbers do not predict how frontier models would score, because nothing was measured against hosted APIs. Prices for hosted models are taken from official pages and marked as illustrative.
2. **GPU-only items were not run.** QLoRA/bitsandbytes, vLLM serving, speculative decoding and a real Kubernetes deployment (kind failed under rootless podman) are documented and their configurations validated, but none of them was executed. The fine-tune in T1B/C2 is a CPU LoRA on a 0.5B model.
3. **Native-speaker review needed.** All Tamil and Hindi text (prompts, test sets, transcripts and examples in M12, C1, E1 and elsewhere) was written without a native-speaker review. Have it reviewed before you rely on it.
4. **Synthetic data and small test sets.** Every dataset is synthetic, and evaluation sets are small (tens of items). Confidence intervals are reported where they matter, and they are wide.

5. **Android.** No device or emulator was available. Kotlin code in E1 is illustrative and was not compiled. On-device throughput was approximated with an Ollama proxy on the VM CPU. E1's memory cap could not be enforced on the box.
6. **Licences.** Some models used in labs carry non-commercial licences (for example MMS-TTS, CC-BY-NC). Check every licence before using a model in a product.
7. **SaaS platforms.** Zapier, Make, hosted tracing and the enterprise agent platforms were not used, because no accounts were created. Their recipes and comparisons come from official documentation.
8. **Portfolio gaps (closed) and two E1 measurement bugs.** Lab 13.4 first found that the capstones lacked demo scripts and that C3 and E1 lacked READMEs, a Dockerfile and (E1) tests. These were closed on 12 October 2026 (M13 §13.3.6). Writing E1's tests exposed two measurement bugs. Re-scored from saved outputs, Whisper-small's Tamil WER is 61.4%, not the 31.7% first printed, and the SmolVLM GSTIN counts are one higher. The scoring code is fixed, and the E1 booklet was rebuilt with the corrected figures and an erratum. The other Tamil figures (tiny, base, noisy audio) cannot be re-scored without re-running the models and are shown as n/a.
9. **Version drift.** Versions and links were verified on 11–12 October 2026 and will drift. Re-verify them (`research/`) before relying on them.

Appendix A. Week 0 — Python bridge for Kotlin/Java developers (optional)

Kotlin / Android idea	Python equivalent	Covered in
`data class`	`@dataclass` / Pydantic `BaseModel`	M1 §1.3
Null safety `String?`	`str | None` + type checker (mypy/pyright)	M1 §1.3
Coroutines, `suspend` , `Dispatchers.IO`	`async def` , `await` , event loop, `asyncio.to_thread`	M1 §1.4
`async { }` + `awaitAll()`	`asyncio.gather` / `TaskGroup`	M1 §1.4
Retrofit + OkHttp interceptors	httpx client + event hooks / tenacity retries	M1 §1.4
Gradle version catalog	`pyproject.toml` + lock file (uv) / pinned `requirements.txt`	M1 §1.2
JUnit	pytest	M1 §1.5
`sealed class`	`Literal` / `Union` types, `match` statement	M1 §1.3

Appendix B. Glossary

The consolidated course glossary is **`99-glossary.pdf`** : every module's glossary merged and de-duplicated, with module cross-references and an acronym list. It is generated by `labs/course-glossary/build_glossary.py` from the module sources and the seed glossary in `course-plan.md` (section 8).